

TripletJEPA — Results Summary

Evaluation protocol (CIFAR-100, frozen encoder): - k-NN@20 on L2-normalized features (cosine distance) - Linear probe: SGD 50 epochs on train features, test accuracy - Corrupt view: 60% block mask (4x4 patches) - SmallResNet ~2.8M params unless backbone=resnet50

Completed runs loaded from results/*/checkpoint_last.pt

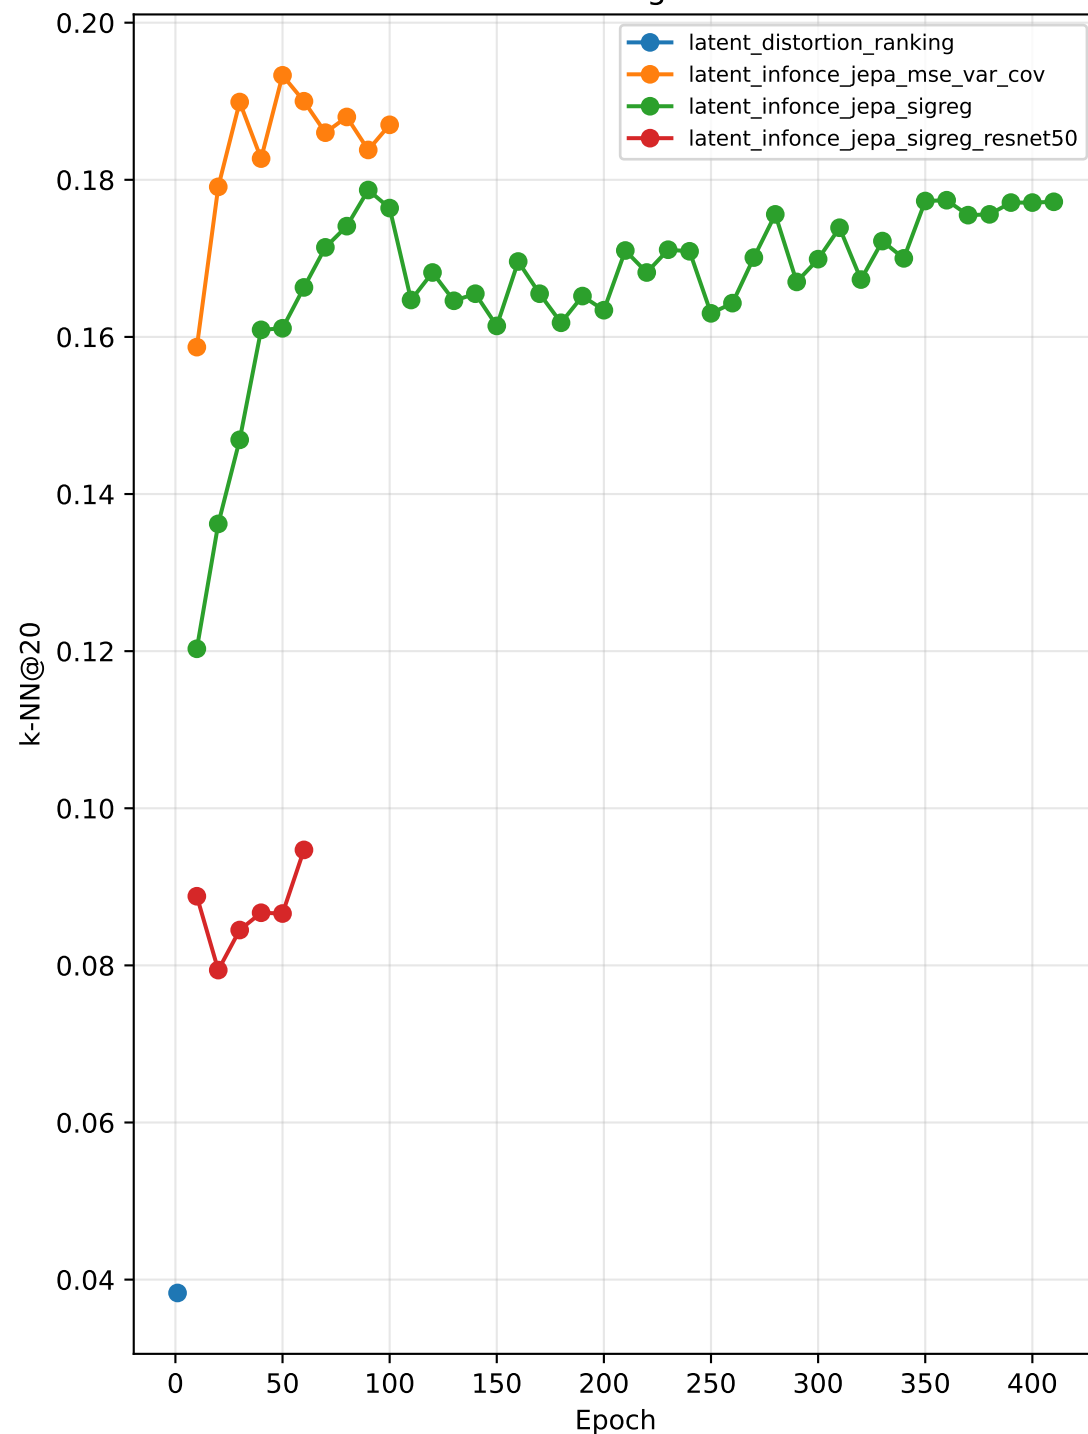
Generated from checkpoint histories. LaTeX source: paper/main.tex

CIFAR-100 SSL Results (frozen encoder eval)

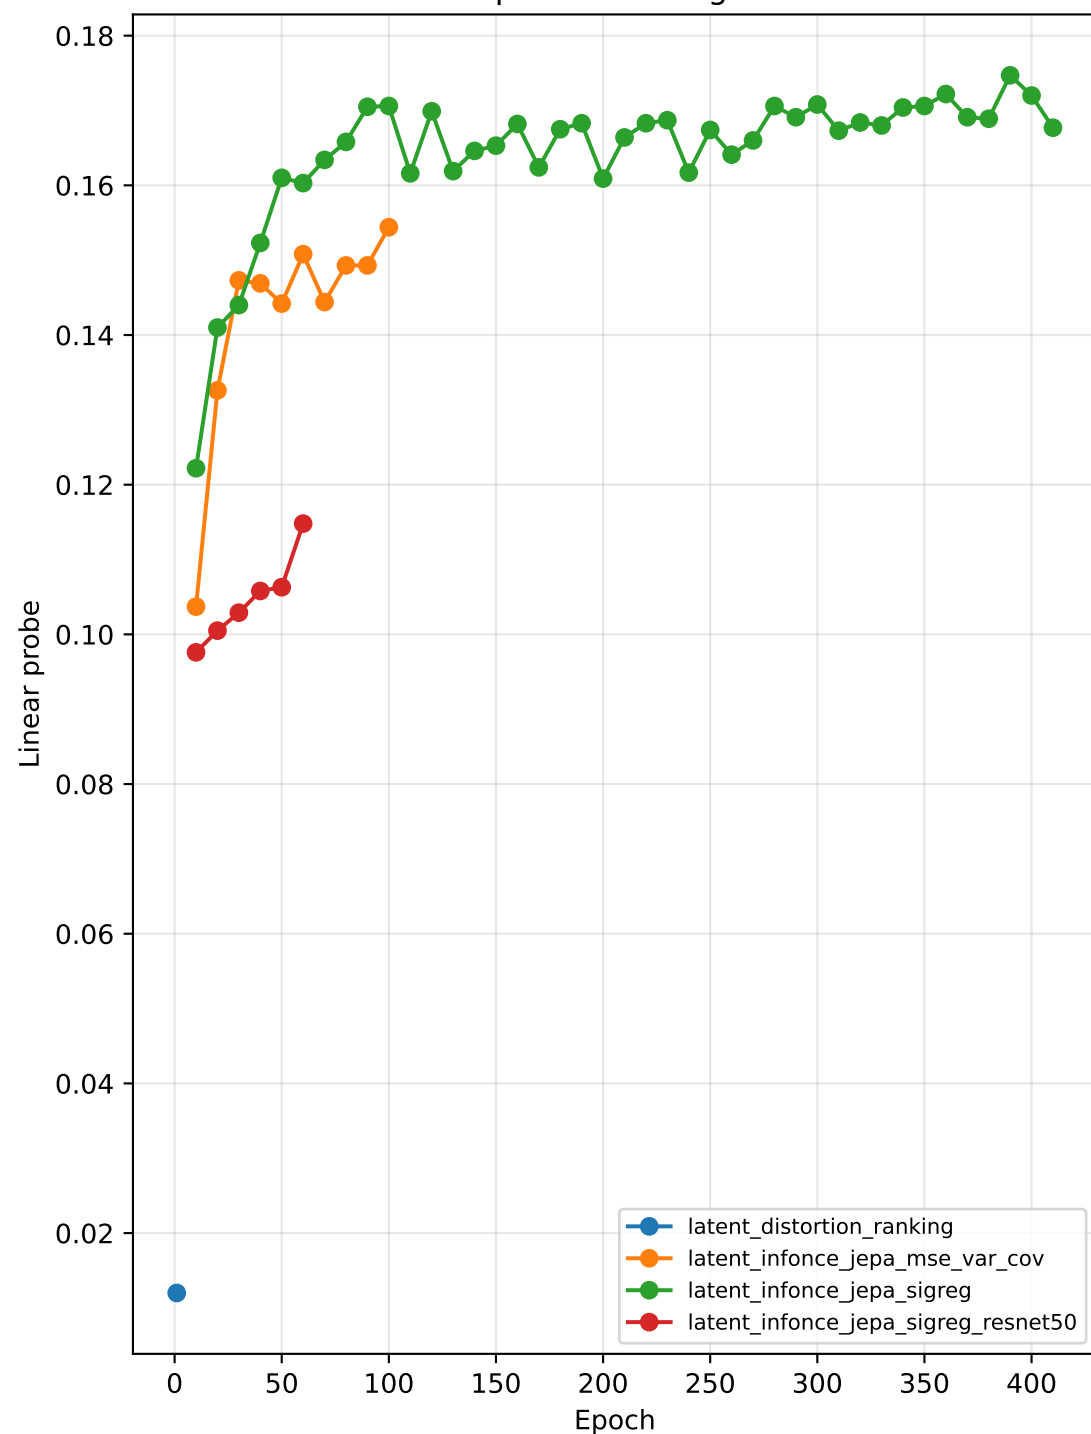
Run	Backbone	Ep	Best k-NN	@	Best Linear	@	Last k-NN	Last Linear
ent_distortion_rank	resnet	1	0.038	1	0.012	1	0.038	0.012
nt_infonce_jepa_ms	resnet	100	0.193	50	0.154	100	0.187	0.154
nt_infonce_jepa_s	resnet	418	0.179	90	0.175	390	0.177	0.168
nt_infonce_jepa_si	resnet50	62	0.095	60	0.115	60	0.095	0.115

Training curves from checkpoints

k-NN learning curves



Linear probe learning curves



k-NN vs Linear trade-off

